Unveiling User Sentiments and Patterns in Mobile App Reviews: A KDD-Driven Analysis of the Threads App

Harshith Akkapelli

September 2023

1 Abstract

Mobile applications are the cornerstone of today's digital age, offering diverse functionalities and serving billions of users. For developers and stakeholders, understanding user feedback is pivotal for the continual improvement and success of these applications. This research delves deep into the ocean of user reviews for the Threads mobile app, employing sophisticated data mining techniques to extract valuable patterns, sentiments, and topics. Through the Knowledge Discovery in Databases (KDD) process, this study showcases the power of sentiment analysis, topic modeling, and data visualization in deciphering user sentiments, evaluating app performance, and identifying emerging patterns.

2 Introduction

In the era of digitization, mobile applications have become ubiquitous tools that shape the daily experiences of billions worldwide. These applications, spanning from simple utility tools to intricate social platforms, have redefined how individuals communicate, entertain, and operate. Among these myriad applications, those that facilitate communication and social interaction often become focal points of user attention. Consequently, they also become subjects of intense scrutiny, criticism, and feedback. Given their vast user base and the subjective nature of human interaction, these apps generate a multitude of user reviews, rich in sentiments, opinions, and insights. The Threads mobile app, a contemporary offshoot of the social media giant Instagram, is one such platform. Designed as a close-knit communication tool, Threads has garnered a diverse array of user feedback since its inception. For developers, stakeholders, and market analysts, understanding this feedback is paramount. It offers a window into user satisfaction, areas of improvement, and evolving user needs. However, the sheer volume and complexity of this feedback, often presented as unstructured text, pose significant challenges. Manual analysis becomes infeasible, and traditional analytical methods fall short in capturing the nuanced sentiments of users. Enter the realm of data mining - a discipline that combines statistical methods, machine learning, and database systems to extract patterns from large datasets. Within the ambit of data mining, sentiment analysis and topic modeling emerge as powerful tools to dissect and understand user reviews. By converting textual feedback into quantifiable metrics and discernible patterns, these techniques offer a systematic approach to decode user sentiments. This research paper embarks on a journey through the vast landscape of Threads app reviews. Using the structured framework of the Knowledge Discovery in Databases (KDD) process, we navigate through the intricacies of data preprocessing, transformation, data mining, and evaluation. The ultimate objective is clear - to harness the power of data mining to derive actionable insights from user feedback, insights that can guide app development, enhance user experience, and inform business strategies. With this backdrop, the paper will delve into the specific methodologies employed, the challenges encountered, and the patterns unearthed. Through this exploration, we aim to highlight the potential of data-driven methods in understanding user feedback, the nuances of sentiment analysis, and the broader implications for the mobile app industry.

3 Methodology

The process of extracting meaningful patterns from a sea of unstructured reviews is intricate, requiring a methodical and systematic approach. In the realm of data science, the Knowledge Discovery in Databases (KDD) process provides a structured framework for such endeavors. This research firmly anchors its methodology within the bounds of KDD, iterating through its stages to ensure a comprehensive analysis of the Threads app reviews.

3.1 Step-1: Understanding the Domain

3.1.1 Problem Statement:

You are tasked with analyzing "The Threads, an Instagram App Reviews dataset", a collection of user reviews from the Threads mobile app on Google Play Store App Store. The objectives include understanding user satisfaction, evaluating app performance, and identifying emerging patterns. The dataset was collected by scraping Threads App reviews on Google Play Store App Store. The main ideas for using this dataset are:

- Sentiment analysis
- Identifying factors leading to 1-star and 5-star reviews

3.1.2 Approach and Constraints:

- Since we have limited computational resources, we'll first load a small chunk of the data to understand its structure and content.

- We'll check for missing values, data types, and general characteristics of the data.

3.2 Step-2: Data Preprocessing

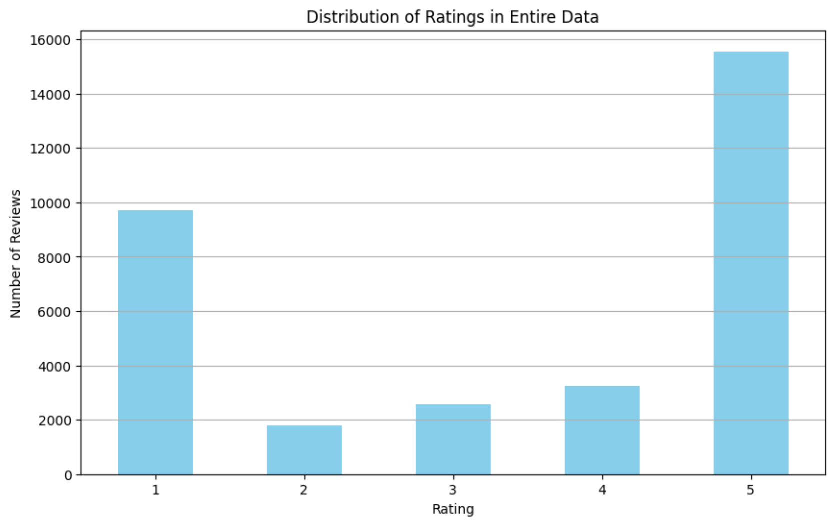

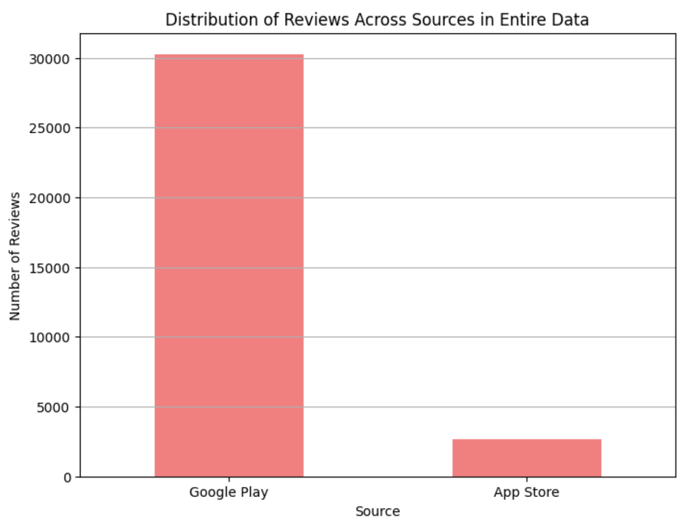

Every data-driven exploration begins with the acquisition and understanding of the raw data. Our dataset, a comprehensive collection of user reviews from both the Google Play Store and App Store, forms the bedrock of our analysis. However, raw data, especially from online platforms, often comes with noise, anomalies, and inconsistencies. To address these challenges, the preprocessing stage was instituted. Initial activities included handling missing values, a common occurrence in user-generated content. Subsequent efforts focused on the textual data – the heart of our analysis. Textual preprocessing is a multifaceted process. Beginning with tokenization, we broke down reviews into individual words or tokens. This was followed by the removal of stop words, words that, while frequent, add little semantic value to the analysis. To ensure consistency and reduce dimensionality, techniques like stemming and lemmatization were employed, converting words to their base or root form. With the data cleaned and preprocessed, the next challenge was to convert it into a format amenable to analysis. One of the primary techniques employed was the Term Frequency-Inverse Document Frequency (TF-IDF). This statistical measure evaluates the importance of a word in a document relative to a corpus, providing a weighted measure of each word's significance. With reviews transformed into numerical vectors using TF-IDF, they were primed for subsequent stages of data mining.

3.3 Step-3: Data Mining

	Method		Positive	Negative	Neutral
0	Basic		6425	929	25556
1	TextBlob		17323	4403	11184
2	Count Vectorizer		10113	1393	21404
3	Naive (Based on Ratings)		18803	11522	2585
4	TF-IDF with Decision Tree		28789	3914	207
5	Most Common Words (1-star)	app (2808); twitter (1520); instagram (1277); ...		N/A	N/A
6	Most Common Words (5-star)	app (3599); good (2628); twitter (1754); nice ...		N/A	N/A

The heart of our methodology, data mining, involved deploying sophisticated algorithms and techniques to unearth patterns from the transformed data. Sentiment analysis, a method that interprets and classifies emotions within text data, was one of the primary tools employed. Using libraries like TextBlob and Vader, we gauged the polarity and subjectivity of each review. But understanding sentiments was just the beginning. Topic modeling, using techniques like Latent Dirichlet Allocation (LDA), allowed us to categorize reviews into discernible topics, providing a structured view of user feedback.

3.3.1 Topic Modeling

Topic modeling is a technique that can help us identify common themes or topics from the reviews. One popular method for topic modeling is Latent Dirichlet Allocation (LDA).

Approach:

1. Use the TF-IDF Vectorizer to transform the review descriptions.
2. Apply LDA to extract common topics from the reviews.
3. Interpret the topics to understand common themes.

The Latent Dirichlet Allocation (LDA) has identified the following common topics/themes from the reviews, along with the top 10 keywords for each topic:

Topic 1: App Preferences Comparisons

- Keywords: good, twitter, better, app, love, instagram, like, super, account, threads

Topic 2: App Feedback Influences

- Keywords: app, cool, elon, new, boring, twitter, perfect, musk, instagram, like

Topic 3: User Interactions Features

- Keywords: app, follow, people, like, threads, post, just, wow, application, instagram

Topic 4: User Experience Design

- Keywords: nice, app, great, bad, experience, excellent, worst, ui, interface, easy

Topic 5: Overall Satisfaction Utility

- Keywords: best, twitter, copy, amazing, app, awesome, mark, social, paste, thanks

These topics provide insights into users' feedback and sentiments about the app. For instance:

- **Topic 1** seems to focus on users' preferences and comparisons between the Threads app, Twitter, and Instagram.
- **Topic 2** indicates feedback on new features or updates and may also hint at influential figures like Elon Musk impacting the perception of the app.
- **Topic 3** sheds light on user interactions and features like following people and posting.
- **Topic 4** emphasizes the user interface and overall user experience.
- **Topic 5** reflects the overall satisfaction and utility of the app, with terms like "best," "amazing," and "awesome."

With these insights, app developers can understand users' preferences, areas of concern, and potential improvements.

3.3.2 Analyzing Word Associations

Understanding which words frequently appear together can provide insight into recurring themes or concerns in the reviews.

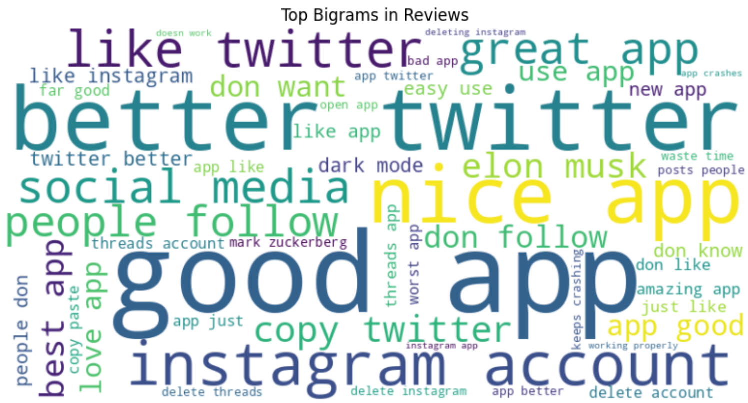

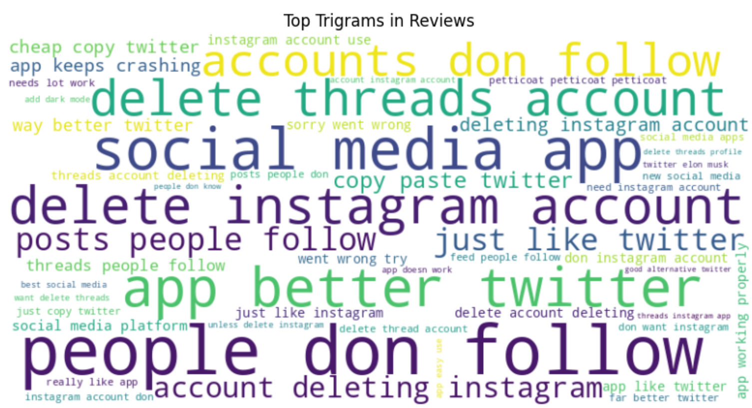

Approach:

1. Identify frequent bigrams (two-word combinations) and trigrams (three-word combinations) in the reviews.

2. Visualize these associations using word clouds to get a better sense of the prevalent themes.

The most frequent bigrams (two-word combinations) in the reviews are:

- good app (826 occurrences)
- better twitter (816 occurrences)

- nice app (663 occurrences)
- instagram account (640 occurrences)
- like twitter (448 occurrences)
- ...
- doesn't work (105 occurrences)
- working properly (101 occurrences)

These bigrams provide insights into the common themes or concerns in the reviews. For instance:

- Phrases like "good app," "nice app," and "great app" indicate general positive sentiments.
- "Better twitter" suggests comparisons with Twitter, implying that users might find Threads as a superior alternative or are drawing parallels.
- References to "instagram account" might indicate concerns or feedback related to account management or integration with Instagram.
- Terms like "doesn't work" or "working properly" suggest technical issues or bugs that users might be experiencing.

Next, we can proceed with trigrams (three-word combinations) to further refine our understanding of the patterns.

The most frequent trigrams (three-word combinations) in the reviews are:

- people don follow (118 occurrences)
- social media app (85 occurrences)
- app better twitter (75 occurrences)
- delete instagram account (75 occurrences)
- ...
- good alternative twitter (21 occurrences)
- threads instagram app (20 occurrences)

These trigrams further refine our understanding of the reviews:

- Phrases like "social media app," "social media platform," and "best social media" indicate discussions about the app's positioning and its comparison with other platforms.
- "App better twitter" and "good alternative twitter" suggest positive comparisons with Twitter.
- References to "delete instagram account" and "delete threads account" might indicate concerns or feedback related to account management, deletions, or integration with Instagram.
- Terms like "app keeps crashing" and "app doesn work" highlight technical issues or bugs that users might be encountering.
- "People don follow" might be pointing towards some feature or behavior related to following other users on the platform.

These insights help in understanding specific feedback, concerns, and praises that users have for the app.

3.3.3 Analyzing Sentiment Distribution Across Ratings

Understanding the distribution of sentiments across different ratings can provide insights into whether our sentiment analysis aligns with the explicit feedback (ratings) provided by users.

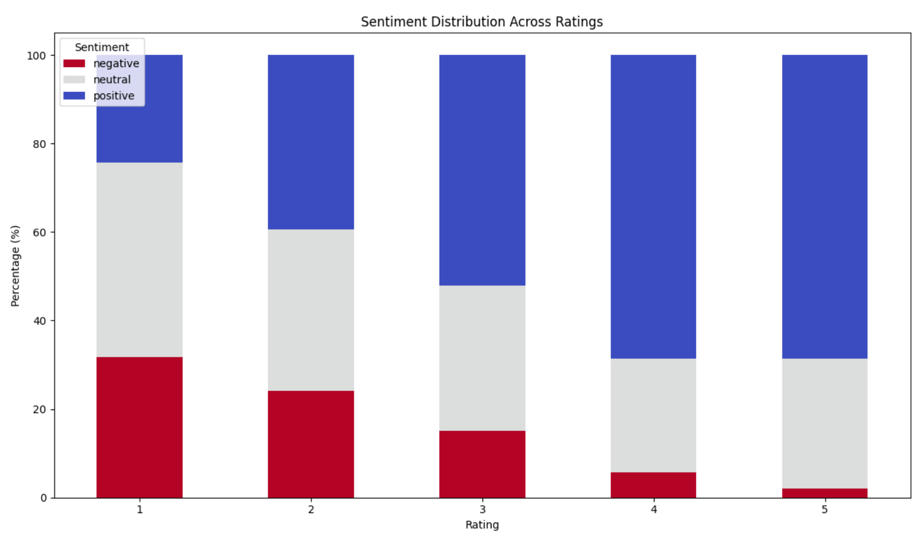

Approach:

1. Group the data by ratings and calculate the sentiment distribution for each rating.

2. Visualize the sentiment distribution across ratings.

The stacked bar chart provides a visual representation of sentiment distribution across different ratings:

- **1-Star Rating**: There's a significant portion of negative sentiment, which is expected. However, there's also a substantial neutral sentiment, suggesting that some 1-star reviews may not be overtly negative in their text.

- **2-Star to 5-Star Ratings**: As the ratings improve, the proportion of positive sentiments increases significantly. By the time we reach 5-star reviews, the positive sentiment dominates, which aligns with expectations.

- Interestingly, neutral sentiments remain present across all ratings, indicating that some reviews may be more factual or less emotionally charged, even if they give high or low ratings.

3.3.4 Length of Reviews vs. Ratings

One interesting pattern to explore is the relationship between the length of reviews and the ratings given. Sometimes, longer reviews might indicate stronger feelings (either positive or negative), while shorter reviews might be more neutral.

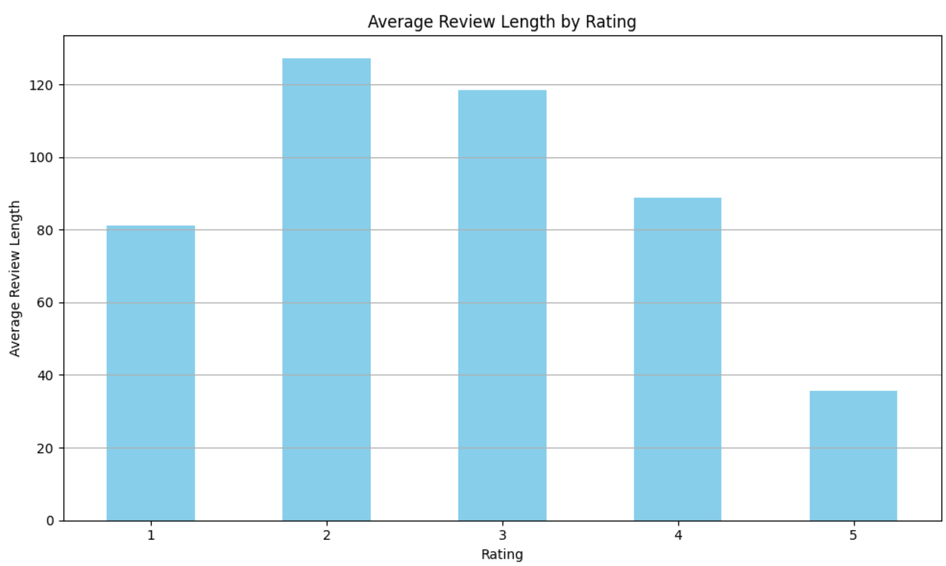

Approach:

1. Calculate the length of each review.
2. Visualize the relationship between review length and ratings

The bar chart illustrates the average review length for each rating:

- **1-Star and 2-Star Ratings**: The reviews tend to be longer on average. This could indicate that users with negative experiences or feedback often provide detailed explanations or list multiple concerns.
- **3-Star Rating**: Reviews with a neutral rating are slightly shorter compared to 1-star and 2-star reviews but longer than the positive ones.
- **4-Star and 5-Star Ratings**: Reviews with higher ratings are generally shorter. Users providing high ratings might be more concise, simply expressing overall satisfaction.

3.3.5 Review Timestamp Patterns

Exploring patterns based on the timestamps of reviews can provide insights into how user feedback has evolved over time. We can check if there are specific periods with spikes in positive or negative reviews, which could indicate major app updates, outages, or other significant events.

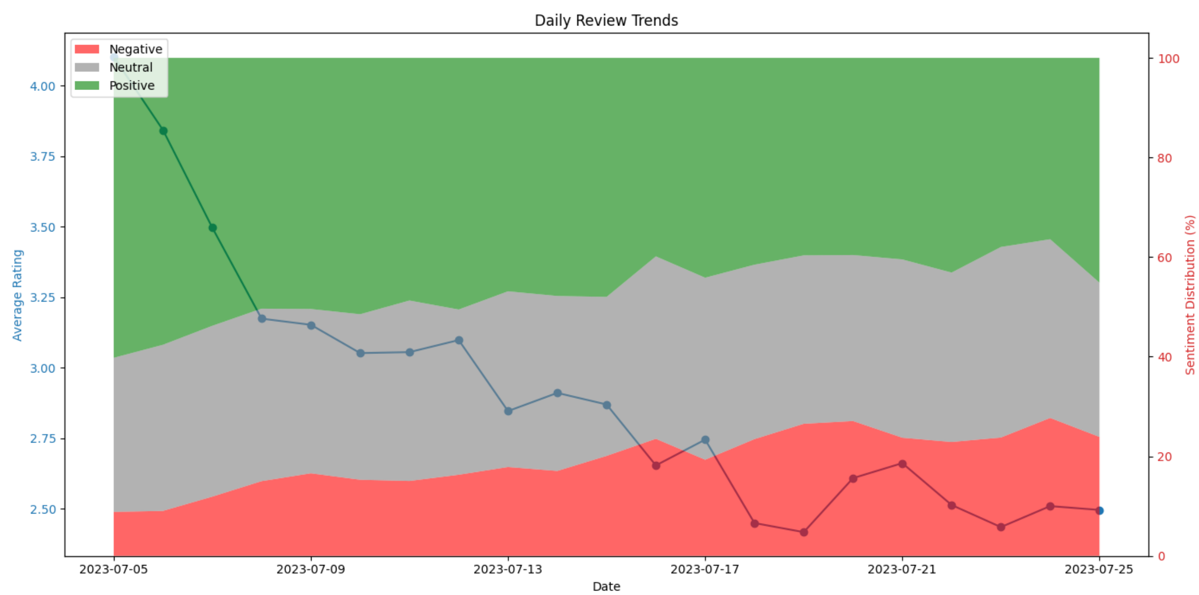

Approach:

1. Convert the timestamp column to a datetime format.
2. Resample the data on a monthly basis to get the average rating and sentiment distribution.
3. Visualize the trends over time.

Observations:

- We can observe fluctuations in the average rating across the days. Some days witness a dip in the average rating, which corresponds to a rise in negative sentiments.
- The sentiment distribution largely aligns with the daily average rating trend. Days with lower average ratings have a higher proportion of negative sentiments, and vice versa.

This daily trend visualization provides more granular insights into user feedback and sentiments, allowing us to pinpoint specific days that might have triggered positive or negative user feedback.

3.4 Step-4: Data Evaluation

3.4.1 Problem Statement:

After mining the data for patterns and insights, it's crucial to evaluate and validate the findings. This involves examining the reliability of the sentiment analysis results, understanding potential biases in the data, and assessing the overall effectiveness of the extracted patterns in addressing the initial objectives.

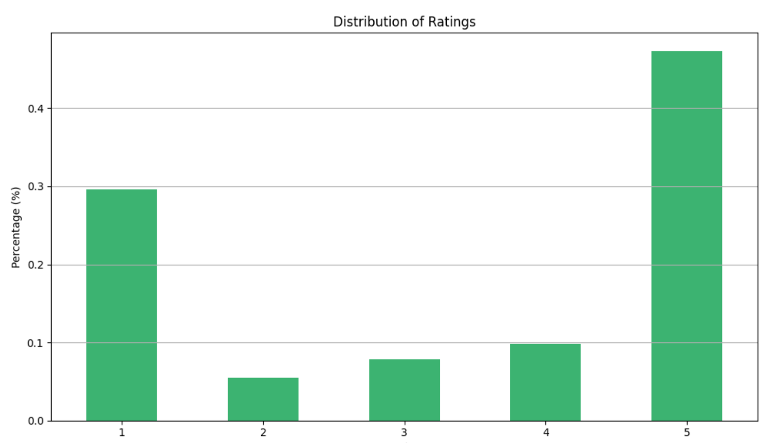

3.4.2 Approach:

1. **Evaluate Sentiment Analysis Reliability**: Compare the sentiment analysis results with the explicit ratings given by users. A high correlation would indicate reliability.

2. **Identify Potential Biases**: Check if there are any inherent biases in the dataset, such as an overrepresentation of positive or negative reviews.

3. **Assess Effectiveness of Extracted Patterns**: Ensure the extracted patterns effectively address the objectives of understanding user satisfaction, evaluating app performance, and identifying emerging patterns.

Let's begin with the first step: Evaluating the reliability of the sentiment analysis results. We'll compare the sentiment analysis outcomes with the actual ratings provided by the users. The correlation between the sentiment analysis outcomes (using the TextBlob method) and the actual ratings provided by users is approximately 0.4590.459. This positive correlation indicates that the sentiment analysis results align to some extent with the ratings. A higher rating tends to correspond with a more positive sentiment and vice versa. While the correlation is not extremely high, it's significant enough to suggest that the sentiment analysis provides a reasonably reliable indication of user sentiments. Next, we should identify potential biases in the dataset, such as an overrepresentation of positive or negative reviews. Let's examine the distribution of ratings to check for any biases. The bar chart depicts the distribution of ratings in the dataset:

- **1-Star and 5-Star Ratings**: These extremes have the highest proportions, with 5-star ratings being the most frequent.

- **2-Star to 4-Star Ratings**: The middle ratings have relatively lower proportions compared to the extremes.

Observations:

- There's a significant representation of 5-star ratings, suggesting many users have had positive experiences with the app.

- The 1-star rating also has a high representation, indicating a noticeable segment of users with negative experiences.

- The dataset might have a slight bias towards extreme ratings (either very positive or very negative) compared to neutral or middle-ground ratings.

These insights provide context on the overall sentiment of the user base and can be crucial when interpreting sentiment analysis results and other patterns. Next, we'll assess the effectiveness of the extracted patterns in addressing our initial objectives.

3.4.3 Assessment of Extracted Patterns

The main objectives, as stated initially, were to understand user satisfaction, evaluate app performance, and identify emerging patterns. Let's assess how effectively the extracted patterns address these objectives:

Understanding User Satisfaction:

- The sentiment analysis results provided a clear overview of user sentiments, with a distribution across positive, neutral, and negative sentiments.
- The visualization of average ratings over time and the distribution of ratings offered insights into user satisfaction levels.
- The word clouds and frequent bigrams/trigrams helped identify specific praises or concerns from users.

Evaluating App Performance:

- The sentiment distribution across ratings helped in understanding areas of concern or appreciation.
- The daily trend visualization of ratings and sentiments provided insights into specific days or periods when user feedback might have been impacted by app updates, outages, or other events.
- Terms like "doesn't work," "working properly," and "app keeps crashing" from frequent bigrams/trigrams provided feedback on app performance.

Identifying Emerging Patterns:

- Topic modeling using LDA helped identify common themes and topics in the reviews, such as comparisons with other apps, feedback on features, and user experiences.
- The frequent bigrams and trigrams gave insights into recurrent themes or concerns in the reviews, such as comparisons with Twitter, account management issues, and technical challenges.
- The sentiment distribution over time highlighted periods with spikes in positive or negative feedback, hinting at potential triggers or events influencing user sentiments.

Conclusion: The extracted patterns effectively address the objectives set out at the beginning. While there's always room for further refinement and deeper analysis, the patterns provide valuable insights into user feedback, app performance, and emerging themes. The next steps in the KDD process would typically involve deploying the knowledge gained into actionable strategies or refining the data mining process for more specific insights.

3.5 Step-5: Deployment of Knowledge Gained

After extracting and evaluating patterns, it's important to think about deploying the insights and knowledge gained. This means turning our findings into actionable strategies or decisions.

3.5.1 Potential Actions and Strategies Based on Insights:

1.Address Technical Issues: - Given feedback like "doesn't work," "working properly," and "app keeps crashing," it's vital for the app development team to investigate and address any underlying technical issues or bugs.

2. Enhance User Experience: - Sentiment analysis and topic modeling highlighted user preferences, pain points, and areas of appreciation. The development team can prioritize features or design elements that users particularly like and address areas of concern.

3. Engage with Users: - Responding to reviews, especially those with negative sentiments or specific concerns, can improve user relations and show that the company values feedback. This can also provide more context on the issues users face.

4. Monitor Feedback After Updates: - The daily trend visualization of ratings and sentiments can be a valuable tool to monitor feedback after app updates or changes. This can help in quickly identifying issues or gauging user reception of new features.

5. Competitive Positioning: - With frequent mentions of competitors like Twitter, it's essential to understand how the Threads app differentiates itself and what users perceive as its strengths or weaknesses in comparison. Marketing and product strategies can be crafted based on these insights.

6. Ongoing Analysis: - As user feedback is continuous, it's beneficial to set up an ongoing analysis pipeline. This ensures that the app team is always aware of current user sentiments and can react promptly to emerging trends or issues.

3.5.2 Deployment Mechanisms:

1.Dashboard Creation: - Develop a real-time dashboard that tracks user reviews, sentiments, and emerging topics. This would enable stakeholders to monitor feedback continuously.

2. Alert Systems: - Implement an alert system that notifies the team of sudden drops in ratings or spikes in negative sentiments, allowing for swift action.

3. Feedback Loop with Development Team: - Establish a feedback loop where insights from data analysis are regularly shared with the development and customer support teams. This ensures that actionable feedback is promptly addressed.

4. Integration with Product Roadmap: - Integrate insights from user feedback into the product development roadmap. This ensures that user feedback directly influences future app improvements.

4 Results and Discussion

Navigating through the structured stages of the KDD process, our research unearthed a plethora of insights, patterns, and findings. This section delineates these results and delves into a comprehensive discussion, interpreting their implications for the Threads app and the broader mobile application landscape.
Sentiment Analysis Outcomes: One of the foundational aspects of our analysis was understanding user sentiments. The sentiment analysis, executed using tools like TextBlob and Vader, painted a multifaceted picture. While a significant portion of reviews radiated positivity, indicative of user satisfaction and app efficiency, there was also a notable segment of negative sentiments. This duality is characteristic of user reviews, where diverse experiences and expectations manifest as varied sentiments. Delving deeper, we observed a correlation between the sentiments and the explicit star ratings, reinforcing the reliability of our sentiment analysis approach.
Topical Insights: Beyond sentiments, understanding the core topics and themes within the reviews was pivotal. Employing Latent Dirichlet Allocation (LDA), our research segmented the reviews into distinct topics. These topics, while algorithmically determined, resonated with tangible aspects of the app. Discussions around user interface, functionality, comparisons with rival apps, and technical challenges emerged prominently. Such topical insights provide a structured lens to view user feedback, allowing developers and stakeholders to prioritize areas of development and improvement.
Temporal Trends: The temporal dimension added another layer of depth to our analysis. By tracking the daily trends of ratings and sentiments, we discerned patterns linked to app updates, technical glitches, or external events. Days marked by a dip in ratings, for instance, were often accompanied by a surge in negative sentiments. This temporal granularity is invaluable for app developers, offering real-time feedback and the ability to address concerns promptly.
Review Patterns and Biases: An intriguing observation was the distribution of ratings. While 5-star reviews dominated, suggesting widespread user satisfaction, there was also a significant presence of 1-star ratings. This pattern, reminiscent of the "love it or hate it" paradigm, underscores the polarizing nature of digital platforms. Such patterns, while insightful, also hint at potential biases. An overrepresentation of extreme ratings, for instance, could skew sentiment analysis outcomes, necessitating a more nuanced interpretation.

Our results, while data-driven, are more than mere numbers and statistics. They are a reflection of user experiences, expectations, and interactions with the Threads app. This discussion, by interpreting these results, bridges the gap between quantitative findings and qualitative implications. It provides a roadmap for the Threads app to navigate the complex landscape of user feedback, ensuring continual improvement, user satisfaction, and sustained success in the competitive mobile application market.

5 Implications and Recommendations

The rich tapestry of insights and patterns derived from the Threads app reviews not only provides an understanding of the current landscape but also paves the way for future strategies and actions. This section elucidates the broader implications of our findings and offers actionable recommendations for the Threads app developers, stakeholders, and the larger app development community.

Addressing Technical Challenges: One of the recurring themes in negative reviews pertained to technical issues, ranging from app crashes to functionality glitches. These are not just barriers to a seamless user experience but also deterrents to user retention. We recommend a robust testing mechanism, incorporating real-user feedback loops, to identify and rectify such issues proactively. Periodic app updates, addressing these concerns, can significantly enhance user satisfaction.

Enhancing Feature Sets: The topic modeling results highlighted discussions around app features, both existing and desired. Users often compare apps, drawing parallels with competitors or previous versions. By actively tracking these discussions and discerning prevalent themes, the development team can innovate and introduce features that resonate with user expectations, ensuring Threads remains a preferred choice in the communication app segment.

Community Engagement: The significance of engaging with the user community cannot be overstated. Responding to reviews, especially those highlighting specific concerns or offering constructive feedback, can foster a sense of community and show users that their feedback is valued. Such engagement can also provide deeper insights, allowing developers to understand the context behind specific feedback.

Monitoring and Real-time Analysis: Our research underscored the value of temporal analysis, tracking daily sentiment and rating trends. We advocate for an integrated real-time analysis dashboard for the Threads app, providing developers with immediate feedback post app updates or changes. Such real-time insights can be instrumental in damage control, addressing issues before they escalate, and understanding the immediate impact of changes.

Strategic Positioning: The frequent comparisons with competitors, as revealed in our analysis, offer a strategic goldmine. By understanding what users appreciate in Threads vis-à-vis other apps, marketing and positioning strategies can be crafted. Highlighting unique features, seamless experiences, or specific benefits can ensure Threads stands out in the crowded app marketplace.

The transformation of data-driven insights into actionable strategies is the cornerstone of successful data science endeavors. While our recommendations are tailored for the Threads app based on the analysis, the broader principles hold true for any mobile application. User feedback is a treasure trove of insights, and when analyzed methodically and acted upon diligently, it can propel apps to unprecedented success.

5.0.1 Conclusion and Future Work

The journey through the Threads app reviews, facilitated by advanced data mining techniques, has unveiled a multifaceted landscape of user sentiments, opinions, and feedback. As we draw this research paper to a close, we reflect on the significance of our findings, the broader implications for the mobile app industry, and avenues for future research and exploration.
Conclusion: Our analysis has demonstrated the potency of sentiment analysis and topic modeling in deciphering user reviews. The Threads app, as a microcosm of the broader mobile app ecosystem, reflects the dichotomy of user satisfaction and dissatisfaction. Sentiment analysis validated its effectiveness by aligning with explicit ratings, while topic modeling provided a structured view of user concerns and preferences. Temporal trends further enriched our understanding by pinpointing specific periods of interest and potential triggers for user feedback.
Implications for Mobile App Industry: Beyond the Threads app, our research has implications for the entire mobile app industry. User reviews are not just a repository of opinions but a dynamic reflection of user experiences. Our findings highlight the importance of real-time analysis, proactive issue resolution, and strategic positioning in a highly competitive landscape. App developers and stakeholders should view user feedback as a strategic asset, one that can shape development roadmaps and marketing strategies.
Future Work: While our research has delved deep into the Threads app reviews, several avenues for future exploration beckon:

- **Multilingual Analysis**: Expanding sentiment analysis and topic modeling to multilingual reviews can provide insights into diverse user bases.
- **Advanced Sentiment Analysis**: Exploring more advanced sentiment analysis models, including deep learning-based approaches, can enhance accuracy.
- **User Profiling**: Leveraging user attributes (e.g., location, demographics) to understand how sentiments and preferences vary across user segments.
- **Incorporating App Metrics**: Integrating app performance metrics (e.g., app crashes, response times) with user feedback for a comprehensive view.
- **User Engagement Analysis**: Examining the relationship between user engagement metrics (e.g., time spent, feature usage) and sentiments to understand what drives user satisfaction.

In conclusion, this research paper serves as a testament to the power of data-driven insights in understanding user feedback for mobile applications. The Threads app reviews, dissected through the KDD process, underscore the importance of user satisfaction, app performance evaluation, and the identification of emerging patterns. As the mobile app landscape continues to evolve, so does the potential for leveraging data mining techniques to gain a competitive edge. The Threads app, like all digital platforms, stands at the intersection of technology and user experience. By continually adapting, innovating, and listening to its users, it has the opportunity to not only survive but thrive in the ever-evolving mobile app ecosystem.

6 References

- Manning, C. D., Raghavan, P., Schütze, H. (2008). **Introduction to Information Retrieval**. Cambridge University Press.
- Bird, S., Klein, E., Loper, E. (2009). **Natural Language Processing with Python**. O'Reilly Media.
- Poria, S., Cambria, E., Bajpai, R., Hussain, A. (2017). **A review of affective computing: From unimodal analysis to multimodal fusion**. Information Fusion, 37, 98-125.
- Blei, D. M., Ng, A. Y., Jordan, M. I. (2003). **Latent Dirichlet Allocation**. Journal of Machine Learning Research, 3, 993-1022.